

PLAGISENSE: AN AI-BASED SEMANTIC PLAGIARISM DETECTION SYSTEM USING SENTENCE-BERT

Sarthak Bagasi¹, Mr. Ankur Jain²

^{1,2}Dr. Akhilesh Das Gupta Institute Of Professional Studies, India.

ABSTRACT

Plagiarism in academic writing has become increasingly difficult to detect due to the availability of digital content and the widespread use of paraphrasing tools. Traditional text-matching systems rely primarily on lexical similarity and therefore fail to identify instances where the text has been reworded but the underlying meaning remains unchanged. This research presents PlagiSense, an AI-powered plagiarism detection system that evaluates semantic similarity using Sentence-BERT embeddings and cosine similarity. The system preprocesses input text through normalization and tokenization, converts sentences into contextual embeddings, and computes similarity scores to determine the degree of semantic overlap. The results indicate that the model can effectively detect rephrased, conceptually similar, or structurally modified content that typical keyword-matching tools fail to identify. The study demonstrates the potential of transformer-based models to improve academic integrity by offering a more reliable method for detecting semantic-level plagiarism.

Keywords: Plagiarism Detection, Semantic Similarity, Sentence-BERT, Transformer Models, Natural Language Processing, Cosine Similarity, Text Embeddings, Academic Integrity.

1. INTRODUCTION

Plagiarism remains a persistent challenge in academic institutions, particularly in an age where digital resources are abundant and easily accessible. Students can conveniently obtain online material and modify it superficially using simple rephrasing tools, resulting in content that appears original but carries the same underlying meaning as the source text. While traditional plagiarism detection systems such as Turnitin and Copyscape have been widely adopted, they primarily depend on lexical overlap and string-matching algorithms. Consequently, they often fail to detect paraphrased or semantically altered content, which limits their effectiveness in ensuring academic honesty.

The emergence of Natural Language Processing (NLP) and deep learning models has made it possible to analyze text at a semantic rather than syntactic level. Transformer-based models, particularly Sentence-BERT (SBERT), have demonstrated remarkable capabilities in capturing contextual meaning through dense vector embeddings. These embeddings enable comparison between text segments based on semantics, offering a more sophisticated approach to identifying plagiarism. This research explores the development and evaluation of PlagiSense, a plagiarism detection system built using SBERT and cosine similarity to identify different forms of plagiarism, including direct copying, paraphrasing, and conceptual similarity.

2. LITERATURE REVIEW

Plagiarism detection has evolved from simple string-matching methods to sophisticated semantic analysis techniques. Early systems relied heavily on lexical similarity, using models such as TF-IDF and cosine similarity to measure word-level overlap. While efficient, these models lacked contextual understanding, making them unsuitable for identifying paraphrased text. Researchers later introduced vector-based models such as Word2Vec and GloVe, which provided distributed word embeddings.

Although these models captured semantic relationships between words, they were insufficient for sentence-level comparisons because they ignored word order and contextual dependencies.

The development of transformer-based models revolutionized NLP, enabling deep contextual representations of text. BERT, introduced by Devlin et al. (2019), provided bidirectional context awareness, but its embeddings were not optimized for semantic similarity tasks. Sentence-BERT, proposed by Reimers and Gurevych (2019), modified BERT using a siamese network structure to generate sentence-level embeddings suitable for semantic comparison. Recent studies have shown that SBERT achieves state-of-the-art performance in tasks such as paraphrase identification, semantic textual similarity, and duplicate question detection. These advancements suggest strong potential for applying transformer-based models to plagiarism detection, particularly in identifying complex paraphrasing and conceptual similarity. This research builds upon these developments by integrating SBERT with cosine similarity to create a lightweight yet efficient plagiarism detection system.

3. METHODOLOGY

The proposed system follows a structured pipeline consisting of text preprocessing, embedding generation, similarity computation, and plagiarism scoring. Input text is normalized by converting to lowercase, removing punctuation and special characters, eliminating stopwords, and tokenizing it into sentences.

Preprocessing ensures consistency and reduces noise in the embedding generation stage. Each sentence is converted into a dense numerical vector using Sentence- BERT. SBERT produces contextual embeddings that represent the semantic meaning of the sentence, enabling deeper comparison beyond mere word matching. Cosine similarity is applied to the embeddings of two sentences. This metric produces a similarity score between 0 and 1, where values closer to 1 indicate high semantic similarity. Sentence-level scores are aggregated to compute the overall similarity between two documents.

The aggregated similarity score is mapped to a plagiarism percentage. Higher similarity values indicate greater likelihood of plagiarism, including paraphrased or conceptually equivalent content. The system is evaluated using a dataset containing original, directly copied, and paraphrased text samples. The evaluation focuses on the model's ability to detect paraphrasing, semantic similarity, and conceptual overlap.

4. RESULTS

The system was tested using various types of text samples. Directly copied text consistently produced similarity scores above 85%, indicating strong semantic alignment. Paraphrased texts, where the wording was changed but the meaning preserved, yielded similarity scores between 60% and 80%. This demonstrates the system's ability to detect deeper semantic connections that traditional tools fail to identify. Unrelated content produced low similarity scores, typically below 40%, confirming that the model does not produce false positives.

Overall, the results show that Sentence-BERT effectively captures semantic meaning and can distinguish between original and plagiarized content with high accuracy. The plagiarism percentages derived from similarity scores provide a clear and interpretable metric for academic use.

5. CONCLUSION

This research presents PlagiSense, an AI-based plagiarism detection system that leverages Sentence- BERT and cosine similarity to identify semantic similarity between texts. The system successfully detects direct copying, paraphrasing, and conceptually similar content, addressing the limitations of traditional keyword-based tools. The experimental results validate the effectiveness of transformer embeddings for semantic comparison, making PlagiSense a reliable tool for promoting academic integrity.

6. FUTURE SCOPE

Although PlagiSense demonstrates strong performance in detecting semantic-level plagiarism, several enhancements can further improve its effectiveness and practical applicability. Future development may focus on integrating multilingual capabilities by incorporating models such as Multilingual BERT or XLM-RoBERTa, enabling the system to process and compare text across different languages, which is essential in diverse academic environments. Additionally, the system can be extended to highlight plagiarized segments within the text, providing users with a visual representation of similarity patterns instead of a single numerical percentage. This feature would make the output more intuitive and beneficial for educators.

Another important direction involves scaling the system to handle large datasets and long-form documents such as theses, dissertations, and research papers. Optimizing computational efficiency through model pruning, quantization, or the use of lightweight transformer architectures would help make the system more responsive and accessible on low-resource machines. Integration with institutional learning management systems (LMS) or online submission portals would further streamline plagiarism checking workflows, allowing automated evaluation of assignments at the point of submission.

Moreover, future enhancements may include the incorporation of more advanced embedding models and domain-specific fine-tuning to increase accuracy in specialized subjects such as law, medicine, or engineering. Incorporating cross-lingual and cross-domain similarity detection can significantly broaden the system's utility.

Finally, with the rise of AI-generated content, the system can be adapted to detect AI-assisted writing patterns and differentiate between human-written and machine-generated text, making PlagiSense more relevant and future-ready. These developments collectively pave the way for a more robust, scalable, and intelligent plagiarism detection framework.

7. REFERENCES

- [1] Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. arXiv:1908.10084.
- [2] Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. NAACL.
- [3] Alzahrani, S. M., Salim, N., & Abraham, A. (2012). Understanding Plagiarism Linguistic Patterns, Textual Features, and Detection Methods. IEEE Transactions on Systems, Man, and Cybernetics.
- [4] Manning, C., Raghavan, P., & Schütze, H. (2008). Introduction to Information Retrieval. Cambridge University Press.
- [5] Mikolov, T. et al. (2013). Efficient Estimation of Word Representations in Vector Space. arXiv:1301.3781.